

Coevolving Memetic Algorithms: A Review and Progress Report

Jim E. Smith

Abstract—Coevolving memetic algorithms are a family of metaheuristic search algorithms in which a rule-based representation of local search (LS) is coadapted alongside candidate solutions within a hybrid evolutionary system. Simple versions of these systems have been shown to outperform other nonadaptive memetic and evolutionary algorithms on a range of problems. This paper presents a rationale for such systems and places them in the context of other recent work on adaptive memetic algorithms. It then proposes a general structure within which a population of LS algorithms can be evolved in tandem with the solutions to which they are applied. Previous research started with a simple self-adaptive system before moving on to more complex models. Results showed that the algorithm was able to discover and exploit certain forms of structure and regularities within the problems. This “metalearning” of problem features provided a means of creating highly scalable algorithms. This work is briefly reviewed to highlight some of the important findings and behaviors exhibited. Based on this analysis, new results are then presented from systems with more flexible representations, which, again, show significant improvements. Finally, the current state of, and future directions for, research in this area is discussed.

Index Terms—Adaptive systems, memetic adaptive evolutionary co-evolutionary.

I. INTRODUCTION

THE PERFORMANCE benefits that can be achieved by hybridizing evolutionary algorithms (EAs) with local search (LS) operators, or the so-called memetic algorithms (MAs), have now been well documented across a wide range of problem domains such as optimization of combinatorial, nonstationary, and multiobjective problems (see [1] for a review, [2] for a collection of recent algorithmic and theoretical work, and [3] for a comprehensive bibliography). In these algorithms, an LS improvement step is commonly performed on each of the products of the generating (recombination and mutation) operators prior to selection for the next population. There are, of course, many variants on this theme; for example, one or more of the generating operators may be null, or the order in which the operators are applied may vary, but these can easily be fitted within a general syntactic framework [1].

In recent years, it has been increasingly recognized that the particular choice of LS operator will have a major impact on the efficacy of the hybridization. Of particular importance is the choice of the move operator, which defines the neighborhood function and, thus, governs the way in which new solutions are generated and tested. Points that are locally optimal with respect to one neighborhood structure may not be with respect to another (unless, of course, they are globally optimal). Therefore, it follows that even if a population only contains local optima, then changing the LS move operator (neighborhood) may provide a means of progression in addition to recombination and mutation. This observation has led a number of authors to investigate and propose mechanisms for choosing between a set of predefined LS operators that may be used during a particular run of a metaheuristic such as an EA (see Section II for further discussion).

Previous papers have reported initial results from a system within which the definitions of LS operators applied within the MA may be changed during the course of optimization. This was named the Coevolving Memetic Algorithm (COMA). It maintains two populations—one of genes encoding for candidate solutions and one of memes encoding for LS operators to be used within the MA. This paper places the COMA framework in the context of algorithmic advances by other authors and then reviews the progression of research and results from the initial simple systems [4]–[6] to the more complex truly coevolutionary systems tested in [7].

In those investigations the emphasis had been placed on evolving the rule-based neighborhood definition but leaving much of the rest of LS algorithm fixed. This restriction is then removed. Results and analysis are presented from a benchmark testing of the developed systems on a commonly used NP-Hard problem: MAX-3SAT. The rest of this paper proceeds as follows.

- Section II draws some parallels between this work and related work in different fields in order to place this work within the more general context of studies into adaptation, development, and learning.
- Section III describes the proposed approach and how it can be used to represent a range of algorithms.
- Section IV summarizes the results and analysis of a set of preliminary experiments using a simple self-adaptive model. These investigate whether, and if so how, the use of adaptive variable-lengths rules benefits optimization.
- Section V reviews experiments considering more general coevolutionary forms of the model. It examines the effect

Manuscript received June 12, 2006; revised June 19, 2006. This paper was recommended by Guest Editor Y. S. Ong.

The author is with the Faculty of Computing, Engineering, and Mathematical Sciences, University of the West of England, BS16 12QY Bristol, U.K. (e-mail: james.smith@uwe.ac.uk).

Digital Object Identifier 10.1109/TSMCB.2006.883273

of different pivot and pairing strategies and noise tolerance within credit assignment and identifies two different modes of behavior that benefit search: discovering and exploiting of regularities in the search space and continuously changing neighborhood definitions.

- Section VI details new experiments in which more of the LS definitions are made open to adaptation. The effects of these changes are assessed in a benchmark comparison on the classic MAX-3SAT problem.
- Finally, Section VII discusses the implications of these results, draws conclusions, and suggests future work.

II. RELATED WORK

The COMA system can be related to a number of different branches of research, all of which offer different perspectives and means of analyzing its behavior. These range from multimemetic algorithms and the self-adaptation of search strategies through coevolutionary, learning, and developmental systems to the evolution of rules in learning classifier systems. The more important ones are briefly outlined in the following sections.

A. MAs With Multiple LS Operators

Although the author is not aware of other algorithms in which the LS operators used by an MA are adapted in this fashion, there are other recent examples of the use of multiple LS operators within evolutionary systems. Ong *et al.* [8] present an excellent recent review of work in the field of what they term “adaptive MAs.” This encompasses the works of Krasnogor on “multimemetic algorithms” [9]–[12] and Ong and Keane, who term their approach “meta-Lamarckian MAs” [13], and hyperheuristics [14]–[17]. In an interesting extension to the use of a set of fixed strategies, Krasnogor and Gustafson have recently proposed a grammar for “self-generating MAs,” which specifies, for instance, when LS takes place [18], [19]. Essentially, all of these approaches maintain a pool of LS operators available to be used by the algorithm and, at each decision point, make a choice of which to apply. There is a clear analogy between these algorithms and the variable neighborhood search algorithm [20], where a heuristic is used to control the order of application of a set of predefined local searchers to a single improving solution. The difference here lies in the population-based nature of MAs so that not only do we have multiple local searchers but also multiple candidate solutions, which makes the task of deciding which LS operator to apply to any given one more complex.

The classification of Ong *et al.* uses the terminology developed elsewhere to describe adaptation of operators and parameters in EAs [21]–[23]. They categorize algorithms according to the way that these decisions are made. One way is to use a fixed strategy; this is termed “static.” Another is to use feedback of which operators have provided the best improvement recently. This is termed “adaptive” and is further subdivided into external, local (to a deme or region of search space), and global (to the population) according to the nature of the knowledge considered. Finally, they note that LS operators may be linked

to candidate solutions (self-adaptive). In this terminology, the COMA algorithm may be local adaptive or self-adaptive.

B. Self-Adaptation and Coevolution

As noted above, the COMA algorithm maintains two populations—one of genes and one of memes. If the populations are of the same size, and selection of the two is tightly coupled, then this instantiation of the COMA framework can be considered as a type of self-adaptation. This will also be referred to as “linked” later on in this paper. The use of the intrinsic evolutionary processes to adapt mutation step sizes has long been used in evolution strategies [24] and evolutionary programming [25]. Similar approaches have been used to self-adapt mutation probabilities [26], [27] and recombination operators [28] in genetic algorithms (GAs) as well as complex generating operators that combined both mutation and recombination [29].

If selection is performed separately for the two populations, then COMA is a cooperative coevolutionary system, where the fitness of the members of the meme population is assigned as some function of the relative improvement that they bring about in the “solution” population. Paredis has examined the coevolution of solutions and their representations [30]. Potter and De Jong have also used cooperative coevolution of partial solutions in situations where an obvious problem decomposition was available [31]. Both reported good results. Bull [32] conducted a series of more general studies on cooperative coevolution using Kauffman’s static NKC model [33]. In [34], Bull and Fogerty examined the evolution of linkage flags in coevolving “symbiotic” systems and showed that the strategies that emerge depend heavily on the extent to which the two populations affect each other’s fitness landscape, with linkage preferred in highly interdependent situations. Bull also examined the effect of different pairing strategies [35], with mixed results, although the NKC systems he investigated used fixed interaction patterns. Parker and Blumenthal’s “punctuated anytime learning with samples” (PAL) [36] is a recent approach to this problem of deciding how to pair members of different populations using periodic sampling to estimate fitness.

There has also been a large body of research into competitive coevolution (see [37] for an overview). Here, the fitnesses assigned to the two populations are directly related to how well individuals perform “against” the other population—what has been termed “predator–prey” interactions. Luke and Spector [38] have proposed a general framework within which populations can coevolve under different pressures of competition and cooperation. This uses speciation both to aid both the preservation of diversity and as a way of tackling the credit assignment problem.

In all the aforementioned cooperative and competitive coevolutionary work, the different populations only affect each other’s perceived fitness, unlike the COMA framework where the meme population can directly affect the genotypes within the solution population. This raises the question of whether the modifications arising from LS should be written back into the genotype (Lamarckian learning) or not (Baldwinian learning). Although the pseudocode and the discussion in the following

sections assume Lamarckian learning, this is not a prerequisite of the COMA framework. However, even if a Baldwinian approach was used, COMA differs from the aforementioned coevolutionary systems because there is a selection phase within the LS; therefore, if all of the neighbors of a point defined by the meme's rule are of inferior fitness, then the point is retained unchanged within the population. If one was to discard this criterion and simply apply the rule (possibly iteratively), the system could be viewed as a type of “developmental learning” akin to the studies in genetic coding (e.g., [39]) and the “developmental genetic programming” of Keller and Banzhaf [40], [41].

III. RULE-BASED MODEL FOR THE ADAPTATION OF MOVE OPERATORS

A. Specifying LS

The LS step can be illustrated by the pseudocode in the first box and has three principal components. The first is the choice of pivot rule, which can either be the “steepest ascent” or the “greedy ascent.” In the former, the “termination condition” is that the entire neighborhood $n(i)$ has been searched ($counter = |n(i)|$), whereas the latter stops as soon as an improvement is found, i.e., the termination condition is ($counter = |n(i)| \vee (best \neq i)$). Note that some authors resort to only considering a randomly drawn sample of size $N \ll |n(i)|$ if the neighborhood is too large to search.

The second component is the “depth” of the LS, i.e., the “iteration condition.” This lies in the continuum between only one improving step being applied ($iterations = 1$) to the search continuing to local optimality [$(counter = |n(i)|) \wedge (best = i)$]. Considerable attention has been paid to studying the effect of changing this parameter within MAs, e.g., [42]. Like the choice of pivot rules, it has been shown to have an effect on the performance of the LS algorithm, both in terms of time taken and quality of solution found.

Local_Search(i):

Begin

/* given a starting solution i */

/* and a neighborhood function n */

set $best = i$;

set $iterations = 0$;

Repeat Until (iteration condition is satisfied)

Do

set $counter = 0$;

Repeat Until (termination condition satisfied)

Do

generate the next neighbor $j \in n(i)$;

$counter = counter + 1$

If [$f(j)$ is better than $f(best)$] **Then**

set $best = j$;

Fi

Od

set $i = best$;

set $iterations = iterations + 1$;

Od

End.

The third and primary factor that affects the behavior of the LS is the choice of neighborhood generating function. This can be thought of as defining a set of points $n(i)$ that can be reached by the application of some move operators to point i . One representation is as a graph $G = (v, e)$, where the set of vertices v are the points in the search space, and the edges relate to applications of the move operator, i.e., $e_{ij} \in G \iff j \in n(i)$. The provision of a scalar fitness value f , which is defined over the search space, means that we can consider the graphs defined by different move operators as “fitness landscapes” [43]. Merz and Freisleben [44] present a number of statistical measures that can be used to characterize fitness landscapes and have been proposed as potential measures of problem difficulty. They show that the choice of move operator can have a dramatic effect on the efficiency and effectiveness of the LS and, hence, of the resultant MA.

B. Adapting the Specification of LS

The aim of this work is to provide a means whereby the definition of the LS operator used within an MA can be varied over time and then to examine whether evolutionary processes can be used to control that variation so that a beneficial adaptation takes place. Accomplishing this aim requires the provision of the following five major components:

- 1) a means of representing an LS operator in a form that can be processed by an EA, i.e., a meme;
- 2) a means for initializing a population of memes and a set of variation operators to be applied to them;
- 3) a means of assigning fitness to the meme population;
- 4) a choice of population structures and sizes as well as selection and replacement methods for managing the meme population;
- 5) a set of experiments, problems, and measurements to permit evaluation and analysis of the system.

The representation chosen for the memes is a tuple $\langle Iterate_Condition, Terminate_Condition, Pairing, Move \rangle$. Once this representation is chosen, this leads naturally to the choice of variation operators to act on that representation within an evolutionary setting and can readily encompass all of the other requirements identified above. The pseudocode in Fig. 1 illustrates the algorithm that was used to undertake the research reported in this paper. Note that although this pseudocode assumes equal-sized populations and synchronous evolution, in general, this need not be the case.

The two conditions in the tuple have been described above and can be easily mapped onto an integer or cardinal representation as desired and manipulated by standard genetic operators. The element *Pairing* indicates how the choice of which members of the two populations to combine for evaluation is managed. The purpose of this element is to allow the system to be varied between two extremes. One is a fully unlinked system, in which, although still interacting, the two populations evolve separately. The other is a fully linked self-adapting system. Here, the memes are effectively extra genetic material inherited and varied along with the problem representation. As is illustrated in the **If ... Else** section of the pseudocode,

```

COevolving Memetic Algorithm for Binary Coded Problems :
Begin
/* given a starting population  $P$  of  $\mu$  solutions and a same-sized population  $M$  of memes */
initialise  $P$  with randomly selected binary genes;
initialise  $M$  with randomly selected memes;
set generations = 0;
set evaluations = 0;
Repeat Until ( run_termination condition is satisfied )
Do

/* Create  $\mu$  offspring in each population */
For child := 1 To child =  $\mu$  Do
/* Create offspring by selection, recombination and mutation, storing the parents id's */
set FirstParent[child] = Select_One.Parent( $P$ );
set SecondParent[child] = Select_One.Parent( $P$ );
set Offspring[child] = Recombine(FirstParent[child], SecondParent[child]);
Mutate(Offspring[child]);

/* Select parents of the new meme */
set Pairing = Get_Pairing( $M$ , child);
If (Pairing = Linked) Then
    set FirstMemeParent[child] = FirstParent[child];
    set SecondMemeParent[child] = SecondParent[child];
Fi
Else If (Pairing = Fitness_Based) Then
    set FirstMemeParent[child] = Select_One.Parent( $M$ );
    set SecondMemeParent[child] = Select_One.Parent( $M$ );
Fi
Else
    set FirstMemeParent[child] = RandInt(1,  $\mu$ );
    set SecondMemeParent[child] = RandInt(1,  $\mu$ );
Esle

/* Create new meme from these parents using recombination and mutation */
set NewMemes[child] = Recombine(FirstMemeParent[child], SecondMemeParent[child]);
Mutate(NewMemesffspring[child]);

/* Finally apply local search to Offspring Using Memes */
set original_fitness = Get_Fitness(Offspring[child]);
/* Applying the rule part of a meme to an offspring produces a set of neighbours */
set Neighbours = Apply_Rule.To.Offspring(Offspring[child], NewMemes[child]);
Evaluate_Fitness(Neighbours);
/* Pivot rule of meme determines choice of neighbour */
set Offspring[child] = Apply_Pivot.Rule.To.Neighbours(Neighbours, NewMemes[child]);
/* Finally update meme fitness according to increase in solution fitness */
set  $\Delta$ fitness = Get_Fitness(Offspring[child]) - original_fitness;
Update_Meme_Fitness(NewMemes[child],  $\Delta$ fitness);
set child = child + 1;
Od
set evaluations = evaluations + 1 + |Neighbours|;
set  $P$  = Offspring;
set  $M$  = NewMemes;
Od
End.

```

Fig. 1. Pseudocode definition of the COMA algorithm.

this range of behaviors can be achieved by encoding the pairing as a value taken from the set $\{Linked, Random, Fitness_Based\}$. This value governs the choice of parents of the meme to be used to perform LS on a given offspring candidate solution.

The representation chosen for the move operators was as *condition:action* pairs, which specify a pattern to be looked for in the problem representation and a different pattern that it should be changed to. Although this representation at first appears very simple, it has the potential to represent highly complex moves via the use of symbols to denote not only single/multiple wild-card characters (in a manner similar to that used for regular expressions in Unix) but also the specifications of repetitions and iterations. Further, permitting the use of different-length patterns in the *condition* and *action* parts of the rule gives scope for the *cut* and *splice* operators working on variable-length solutions.

IV. INITIAL PROOF OF CONCEPT: ANALYSIS OF DIFFERENT FIXED AND VARYING SIZED RULES

A. Initial Implementation

Initial experiments were restricted to considering a simple system and examining its behavior on a well-understood set of test problems. They were designed to find out whether the system was able to evolve useful rules when the rule length naturally matched the structure of the problem and, furthermore, whether the system was able to adapt to an appropriate rule length for different problems. For this reason, it was decided to avoid the various issues concerning population management, pairing strategies, and credit assignment and instead work with a single improvement step, a fully linked self-adaptive system, and a greedy pivot rule. These choices were coded into the chromosomes at initialization, and variation operators were not used on them.

The initial systems only used rules where the *condition* and *action* patterns were of equal length and were composed of values taken from the set of permissible allele values of the problem representation, augmented by a “don’t care” symbol (i.e., #), which is allowed to appear in the *condition* part of the rule but not in the *action* part. The neighborhood of a point i then consists of i itself, plus all those points where the substring denoted by *condition* appears in the representation of i and is replaced by the *action*.

To give an example, a rule $1\#0 \rightarrow 111$ matches the binary string 1100111000 in the first, second, sixth, and seventh positions, and the neighborhood is the set {1100111000, 1110111000, 1111111000, 1100111100, 1100111110}. Note that the string is not treated as toroidal, and the neighbors are evaluated in a random order so as not to introduce positional bias into the LS when greedy ascent is used.

In practice, each rule was augmented by a value *rule_length*, which specifies the number of positions in the pattern string to consider. This permitted not only the examination of the effects of different fixed rule sizes but also the ability to adapt via the action of mutation operators on this value. This representation for the rules allows the use of “standard” genetic operators (uniform/one point crossover, point mutation) to vary this part of the memes’ chromosomes.

B. Test Suite and Experimental Setup

In order to examine the behavior of the system, it was used with a set of variants of test functions whose properties are well known. The first of these was a 64-bit problem composed of 16 subproblems of Deb’s four-bit fully deceptive function [45]. The fitness of each subproblem i is given by its unitation $u(i)$, which is the number of bits set to “1,” i.e.,

$$f(i) = \begin{cases} 0.6 - 0.2 \cdot u(i), & u(i) < 4 \\ 1, & u(i) = 4. \end{cases} \quad (1)$$

In addition to this “concatenated” version (i.e., 4-Trap), a second “distributed” version (i.e., Dist-Trap) was used in which the subproblems were interleaved, i.e., subproblem i was composed of the genes $i, i + 16, i + 32$, and $i + 48$. This separation ensured that in a single application, even the longest rules allowed in these experiments would be unable to alter more than one element in any of the subfunctions. A third variant of this problem (i.e., Shifted-Trap) was designed to be more “difficult” than the first for the COMA algorithm by making patterns that were optimal in one subproblem but suboptimal in all others. Since unitation is simply the Hamming distance from the all-zero string, each subproblem can be translated by replacing $u(i)$ with the Hamming distance from an arbitrary four-bit string. There were 16 subproblems; hence, the binary coding of each ones’ index was used as basis for its fitness calculation.

A generational GA with deterministic binary tournament selection for parents and no elitism was used. One point crossover (with a probability of 0.7) and bit-flipping mutation (with a bitwise probability of 0.01) were used on the problem

representation. These choices were taken as “standard” from the literature, and no attempt was made to tune them to the particular problems at hand. Mutation was applied to the rules with a probability of 0.0625 of selecting a new allele value in each locus (the inverse of the maximum rule length allowed to the adaptive version).

For each problem, 20 runs were made for each population size {100, 250, 500}. Each run was continued until the global optimum was reached, subject to a maximum of one million evaluations. Two performance metrics were considered, namely: 1) the success rate (SR), which is the number of runs finding the global optimum, and 2) the average evaluations to success (AES), which is the mean time taken to locate the global optimum on successful runs. The reason for the large cutoff value was to try and avoid skewing results as what can happen with an arbitrarily chosen lower cutoff, rather than to be indicative of the amount of time available for a “real-world” problem. Note that since one iteration of an LS may involve several evaluations, this allows more generations to the GA, i.e., algorithms are compared strictly on the basis of the number of calls to the evaluation function.

The algorithms used and the abbreviations that will be used to refer to them hereafter are given as follows:

- 1) a “vanilla” GA with no LS;
- 2) a simple memetic algorithm (SMA) using a bit-flipping neighborhood, with one iteration of greedy ascent;
- 3) a version of COMA using a random rule in each application, i.e., with the learning disabled (RandCOMA);
- 4) variants of COMA using rules of fixed lengths in the range [1,10] (1 – COMA, ..., 10 – COMA);
- 5) an adaptive version of COMA (A-COMA), in which the rule lengths are randomly initialized in the range [1, 16]. During mutation, a value of ± 1 is randomly added with a probability of 0.0625, subject to staying in range.

C. Results on the “4-Trap” Function

The results obtained showed that the GA, SMA, and 1-COMA algorithms frequently failed to find the optimum unlike the other COMA variants, which always did. On these problems, there was a clear benefit to using adaptive neighborhood LS, although since the RandCOMA algorithm found the optimum on every run, the SR metric did not provide conclusive evidence that learning was taking place.

Considering the AES, the GA, SMA, and 1-COMA algorithms took longer to locate the optimum, particularly for the smaller population sizes. For all population sizes, there was greater variance in the performance of these three algorithms. Applying Tamhane’s T2 test pairwise to the AES results showed that the performance of the GA, SMA, and 1-COMA algorithms was significantly worse than the rest with 95% confidence for a population of 100 or 250. For a population of 500, 2-COMA joined the significantly slower group.

In short, it could be observed that for fixed rule lengths between 3 and 9, and for the adaptive version, the COMA system-derived performance benefits from evolving LS rules according to both metrics on this function.

D. Results on Variants of Trap Function

For the “Shifted-Trap” function, the performances of the GA and SMA were not significantly different from those on the unshifted version. This was because these algorithms solved the subproblems independently and so were “blind” to whether the optimal string for each was different.

The COMA results exhibited the same pattern of behavior noted above, i.e., fast reliable problem solving for all but 1-COMA and 2-COMA, and even for these two, the AES results were statistically significantly better than GA or SMA.

Considering Dist-Trap, the GA, SMA, and Rand-COMA failed to solve the function to optimality in any run, regardless of population size. While the SR for COMA was less than for the other problems (typically 10–15/20 for a population size of 100 and 15–20/20 for a population size of 250), the same pattern was observed, with better performance (SR and AES), for the adaptive version and fixed rule lengths in the range of 3–5, tailing off at the extremes of the length range.

E. Evolution of Rule Base

The results discussed above are promising from the point of view of improved optimization performance but require some analysis and explanation. The deceptive functions used were specifically chosen because the GA theory suggests that they are best solved by finding and mixing optimal solutions to subproblems. Thus, the GA failed to solve the function when the crossover operator was not suited to the representation (Dist-Trap).

Considering the action of a single bit-flipping LS operator on these “trap” subproblems, a search of the Hamming neighborhood of a solution will always lead toward the suboptimal solution when the unitation is 0, 1, or 2 regardless of pivot rule. Additionally, the greedy search of the neighborhood will lead toward the deceptive optimum 75% of the time when the unitation is 3. This explains the poor results of the SMA and 1-COMA algorithms.

The behavior of the A-COMA algorithm was examined by plotting the population mean against the time of the rule length, the specificity of the condition (the proportion of values set to #), and the unitation of the action.

For the 4-Trap function, the system rapidly evolved medium-length (3–4) general (specificity < 50%) rules whose action was to set all the bits to 1 (mean unitation 100%). Closer inspection of the evolving rule base confirmed that the optimal subproblem string was being learned and applied.

For the Shifted-Trap function, where the optimal subblocks are all different, the rule length decreased more slowly from its initial mean value of 8. The specificity also remained higher, and the mean unitation remained at 50%, indicating that different rules were being maintained. This was borne out by closer examination of the evolved rule sets.

The behavior on Dist-Trap was similar to that on 4-Trap, albeit over a longer timescale. The algorithm could not possibly be learning specific rules about subproblems, since no rule was able to affect more than one locus of any subproblem. Rather, the system learned the general rule of setting all bits to 1. The rules were generally shorter than for 4-Trap, which means

that the number of potential neighbors was higher for any given rule. The high incidence of #’s meant that the rule length defined a maximum radius in Hamming space for the neighborhood, rather than a fixed distance from the original solution. These two observations, together with the longer times to solution, suggest that when the system was unable to find a single rule that matched the problems’ structure, a more diverse search took place using a more complex neighborhood, which slowly adapted itself to the state of the current population of solutions. Full details of these experiments and analysis may be found in [4].

F. Benchmark Test: Protein Structure Prediction (PSP)

Although these results were promising, the test problems used deliberately leaned heavily toward the “building block hypothesis” school of thought in their design. Accordingly, in [5] and [6], the A-COMA algorithm was benchmarked against the standard GA and SMA on a well-known combinatorial optimization problem. The PSP problem concerns the prediction of the “native” three-dimensional form of a protein from knowledge of the sequence of its constituent amino acid residues. Dill’s HP model [46] provides an estimate of the free energy of a fold of a given instance based on the summation of pairwise interactions between the amino acid residues. It is a “virtual residue” model, that is to say that each amino acid residue is modeled by a single atom, whose properties are reduced to a quality of being hydrophobic (H) or hydrophilic (P). Thus, a sequence of l amino acid residues is represented by $s \in \{H, P\}^l$, and the space of valid conformations is restricted to self-avoiding paths on a selected lattice, with each amino acid located on a vertex. Hydrophobic units that are adjacent in the lattice but nonadjacent in the sequence add a constant negative factor to the energy level. All other interactions are ignored, and a fixed penalty is added to infeasible folds. Despite its apparent simplicity, finding the global minimum of the HP model for a given sequence has been shown to be NP-complete on various lattices [47], and EAs have been widely applied since [48].

Twenty instances and parameter settings were taken from [10], which use a two-dimensional (2-D) triangular lattice. For each combination of algorithm and instance, 25 runs were made; each continued until the global optimum was reached, subject to a maximum of one million evaluations. Following [49], the representation used a *relative* encoding where alleles come from the set $\{leftback, leftforward, front, rightforward, rightback\}$ and represent the direction of the next move on the lattice from the point of view of the head of the growing chain. The generational GA used the (500 + 500) selection. One point crossover was applied with a probability of 0.8, and a double mutation (which has the effect of causing the mutation point to act as a pivot) was made with a probability of 0.3. All other settings for mutation probabilities were the same as above.

In addition to the GA and the SMA, we tested versions of COMA using a randomly created rules (i.e., with the learning disabled) and steepest (SRand) or greedy (GRand) pivot rules and fully linked adaptive versions of COMA with the two pivot rules (i.e., CLS and CLG). These results are analyzed according

to effectiveness (SR) and efficiency (AES) as before, plus the mean best value found [i.e., mean best fitness (MBF)].

G. Results on PSP

From a total of 500 runs, the SRs in descending order were CLS (337), SMA (202), CLG (127), GRandom (120), GA (83), and SRand (44). A nonparametric Friedman's test for k -related variables shows that the differences in SR between algorithms was significant. A series of paired t -tests confirmed that the results for the CLS algorithm were significantly better than any of the others with over 95% confidence. This difference was particularly noticeable on the longer instances. Of the other results, the SMA performed well on the shorter instances, and the CLG and GRandom results were surprisingly similar, possibly because the noise inherent in the greedy ascent mechanism created problems for CLG's credit assignment mechanism. Significantly, whatever the form of the LS, all but one of the MAs perform much better than the simple GA.

The least successful algorithm was SRand, and even when it was successful, it took far longer than all of the other algorithms. Like GRandom, it used randomly created rules to define the neighborhood for each solution in each generation. However, unlike GRandom, it searched the whole of each neighborhood. Since many of the random rules would be short or have quite low specificity, this causes large neighborhoods and a consequent increase in the AES values for the same number of evolved candidate solutions considered. It is possible that left to run for longer, the SR of the SRand would have been improved.

Among them, the GA was always the fastest, followed by the SMA. The greedy COMA algorithms were faster than their steepest ascent counterparts. A two-way analysis of variance (ANOVA) showed that both instance and algorithm were significant factors. Post hoc analysis using Tamhane's T2 test confirmed the ordering $GA < \{SMA, CLG\} \leq \{GRand, CLS\} < SRand$, with 95% confidence where $<$ is used and 90% confidence where \leq is used between groups. Ordering within groups is achieved by changing values, but these were not significant at these confidence levels.

Most algorithms found the global optimum for the shorter instances. Therefore, when comparing performance on the basis of the quality of the best solutions found, i.e., MBF, only results for the longer and harder instances were considered. It was observed that CLS reached consistently higher values and with a smaller variance in performance than the others and that the SRand algorithm was correspondingly worse. A two-way ANOVA test on the values for the best solution found in each run, with instance number and algorithm as the factors, confirmed the significance of the algorithm in determining the performance. Post hoc testing using Tamhane's T2 test at the 95% confidence level gave the groupings $CLS > \{CLG, SMA, GRand\} > GA > SRand$.

H. Analysis of Meme Evolution on PSP

A number of test runs of CLS were made in which the evolving memes were saved at regular intervals. These showed a strong tendency toward short rules of the form $\{### \rightarrow leftback rightback\}$ or $\{### \rightarrow leftback leftforward\}$. On

a 2-D triangular lattice, both of these rules act to bring residues i and $i + 2$ into contact, and these patterns could be thought of as the 2-D equivalent of representing a single turn of an alpha helix. Experimentation on a square 2-D lattice showed a tendency toward rules of the form $\{### \rightarrow lll\}$ or $\{### \rightarrow rrr\}$, which is also the shortest path that brings two residues into contact.

The use of the word "tendency" should be noted here. In most cases the rule set continued to contain a number of different rules of varying lengths. It has been argued above that in addition to the extra scalability attained by identifying and reapplying regular structural motifs, the presence of a diverse evolving rule set means that the neighborhood structure defining which points around the current population are examined is continuously changing. Thus, even if the population is converged to a single point, which is locally optimal according to most neighborhood structures, eventually, a rule may evolve for which the neighborhood of that point contains a fitter solution. This can be thought of as continually testing new search landscapes to look for "escape routes" from local optima.

Looking back to the results for the GRandom algorithm, in which the rules defining neighborhoods were created at random, this "changing landscape" effect was noticeable in the superior SRs to the SMA. The fact that the CLS algorithm was the best performer according to both SR and MBF metrics points to both modes of operation having a positive effect.

V. EXTENSION TO TRUE COEVOLUTION

Having established the basic principle of evolving memes, which coded for LS rules as a means of enhancing optimization performance in MAs, the next series of experiments used a full coevolutionary model. The aims were to explore the effects of different pivot rules and linkage strategies and to test the hypothesis of the different modes of operation suggested above. Again, the benchmarks for comparison used a GA and an SMA employing a bit-flipping greedy hill climber. COMA with adaptive rule lengths was used with all possible combinations of greedy or steepest ascent with linked, random, or fitness-based pairing strategies. In the latter case, memes were selected to be used via binary tournaments. The fitness of each meme taken to be $\Delta fitness$, i.e., the fitness improvement caused when it was last applied to a candidate solution. The implications of this credit assignment strategy will be discussed later. These algorithms are referred to as CXY, where X comes from the set $\{L(inked), R(andom), T(ournament)\}$, and Y is one of G (Greedy) or S (Steepest). In order to tease out statistical significance on the harder functions used here, 50 runs were made of each algorithm on each problem instance, with a population size of 500. All other experimental settings were the same.

A. Exploiting Search Space Regularities Gives Scalability

The results summarized above suggested that COMA was able identify and utilize regularities in the problem space. This was investigated further using multiple length variants of two well-understood test functions. The first of these comprised

multiple concatenated copies of (1), with lengths in the range $\{40, 60, 80, \dots, 200\}$.

As expected, the results for SMA were extremely poor. The next worse algorithm was CRS. Out of the 50 runs for each size, the SR steadily decreased 50 at length 40 to 5 at length 100 and zero above that. All the other algorithms showed SR of 49 or 50 up to length 160, but only CLS (39), CLG (50), and CTG (50) solved the 200-bit problem.

The AES results were revealing. The GA was faster than Grand and SRand, but the increase in AES with length was worse than linear. The AES results of the successful COMA variants and analysis of the evolving rule bases supported the hypothesis of discovering and exploiting regularities. In this case, it meant identifying a rule that gives the optimal solution to the subproblems and then applying it to each subproblem in the string in successive generations. CLG was the fastest algorithm, followed by CTS and CLS, and all three were near-linear. For example, a linear regression of AES to length for CLG fitted the data with a correlation coefficient of 0.97.

The second test function was Watson's hierarchical if-and-only-if (H-IFF) function, which is a highly epistatic problem designed to examine the virtues of recombination. At the bottom level, fitness is awarded to matching pairs of adjacent bits in a solution s , i.e.,

$$f_1 s = \sum_{i=0}^{l/2-1} 1 - \text{XOR}(s_{2i}, s_{2i+1}) \quad (2)$$

and this process is applied recursively so that a problem of size $l = 2^k$ has k levels. In each ascending level, the number of blocks is reduced by a factor of 2, and the fitness awarded for each matching pair is increased by a constant factor, which in our case is 2. This problem has a number of Hamming suboptima and two global optima corresponding to the $u(i) \in \{0, 1\}$. Problem sizes $l \in \{8, 16, \dots, 512\}$ were used, corresponding to three to nine levels. Note that for $l > 16$, the length of the blocks to be identified and matched at the highest levels far exceeded the maximum rule length.

All of the MAs had higher SRs than the GA, and again, the same three variants of COMA exhibited the best performance. For example, out of 50 runs with $l = 128$, the SR values were 0 (GA, CRG, CRS), 2 (CTG), 4 (MA), 38 (CLG), 43 (CLS), and 49 (CTS). Only the CLG (10) and CTS (11) variants solved the 256-bit problem. Considering different pivot rules, the same pattern was observed as on other problems: the greedy ascent versions found the optimum faster (lower AES) than the equivalent steepest ascent versions but not as reliable (lower SR). An ANOVA on the MBF results confirmed that the performance was statistically significantly different with 95% confidence. Post hoc analysis showed that the CLG, CLS, and CTS variants had a higher MBF than all other algorithms but did not significantly differ.

Analysis of the evolving rule base suggested that for the H-IFF problem, the improved scalability arose from the system to making a decision between the "1's" and "0's" blocks and then applying this throughout the string. The choice between 1's and 0's appeared to occur with equal probability.

B. Escaping Local Optima by Changing Neighborhoods

The performance improvements exhibited above clearly arise from a situation in which the adjacent epistatic interactions within the problem give rise to patterns in the search space that can be exploited by COMA. The two 64-bit variants of the 4-Trap function described above were used to examine the behavior when this is not the case. As reported above, the SMA and GA completely failed to solve the Dist-Trap functions, and this situation did not change by allowing more runs. The only algorithms that ever located the global optimum were the CLG, CLS, and CTS; all three of which always did so. The AES ranking was $\text{CLG} < \text{CTS} < \text{CLS}$; differences are significant with 99.9% confidence.

On the Shifted-Trap function, the SMA found the optimum 45 times out of the 50 runs, and the random-pairing steepest ascent version of COMA (CRS) 39 times. The GA and all other variants of COMA always located the global optimum in the time allowed. It might be expected that attempting to reuse a pattern on different subproblems would hinder the progress of the COMA algorithms. However, in fact, the mean solution time was not significantly different to that of the GA for all but SMA and the CRS variant, and there was a noticeable reduction in the variability of time to solution.

C. Choice of Pairing and Pivot Strategies in Coevolution

The results presented above showed that the choice of pivot and pairing strategies is crucial and intertwined.

As expected, the greedy variants almost always used less evaluations than the steepest ascent equivalents on successful runs. One obvious explanation for this is that when a "good rule" (e.g., $#### \rightarrow 1111$ for 4-Trap) is applied to a candidate solution, it will often match in a number of places and cause the same fitness improvement in several. Unlike memes with a greedy pivot rule, which stop evaluating neighbors as soon as the first improvement is found, those with a steepest pivot rule will evaluate the whole neighborhood. Thus, the number of evaluations used to achieve the same improvement and, hence, the AES will be higher.

In the case of random pairing, there is no selective pressure in the meme population, so the rule base will remain diverse until genetic drift causes an eventual convergence. For the Shifted-Trap function, this was not a problem, since it is desirable to maintain different *condition* parts of the rules for different subproblems, and CRG performed well. However, for CRS, and CRG on the other problems, it prevented identification and use of appropriate rules, and a corresponding decrease in performance was observed.

In the linked variants, the selective pressure toward the evolution of good rules is created implicitly via a continued association with fit solutions, and CLG outperformed CLS on the binary problems (but not PSP).

By contrast, for some problems (but not all), the extra noise introduced by using a greedy ascent was sufficient to "fool" the simple credit assignment mechanism used in these experiments. Thus, a good rule will only get a low fitness if the first match only leads to a small improvement, whereas larger improvement

(and, hence, fitness) might be seen if it was applied elsewhere in the solution. Another source of noise is the choice of partner. It is possible that a more sophisticated method such as Paredis's Life Time Fitness Evaluation (in which a running average of the last 20 pairings is used) or Parker and Blumenthal's PAL with Sampling may well provide a more stable and robust credit assignment mechanism while retaining the speed benefits of a greedy ascent pivot rule. However, these would not be simple to use in our situation with memes evolving at the same rate as the solutions, unless the selection pressures were much greater.

VI. EXTENDING ADAPTABILITY AND BENCHMARK COMPARISON 2: MAX-3SAT

A. Problem Description

The MAX-SAT problem is a classical combinatorial optimization problem consisting of a number of Boolean variables and a set of clauses built from those variables. A full description and many examples can be found in [50]. For each length {50, 100, 250}, the first 25 were taken from the sets of uniformly randomly created satisfiable instances around the phase transition (in terms of hardness) where there are approximately 4.3 clauses per variable. The same experimental setup was used as for the previous experiments, with a maximum of 500 000 evaluations per run. Each algorithm was run ten times on each instance, giving 250 runs for each combination of algorithm and length.

Like the GA and SMA with steepest (SMA-S) and greedy (SMA-G) ascent, COMA variants with linked or fitness-based pairing were used with fixed pivot strategies as before and also with the pivot condition in the rules randomly initialized and subject to mutation at a rate of 0.0625 (CLA/CTA). In addition, each algorithm was run with and without the ability to use a # symbol in the action part of the rule, which was taken to denote that the symbol in the solution should be inverted. This is denoted as CXY-I and provides the ability to create the rule $\# \rightarrow \#$, which is equivalent to an SMA.

It should be noted at the outset that the random way in which these instances are created does not provide any structural regularities for the COMA algorithms to exploit, so the second "changing neighborhood" *modus operandi* might be expected.

B. SR

Table I shows the number of success from 250 runs. No algorithm found a solution for the problems with 250 variables, so it is perhaps worth reiterating that no time was spent tuning the underlying GA.

As can be seen, for the 50 variable instances, the SMAs have the highest SRs, followed by the variants of COMA with linked rule pairing. For the longer instances, all methods are much less successful, and many instances are not solved by any algorithms. SMA-G and CLG-I show the same performance.

For the shorter instances the steepest ascent strategy is better, but for the longer instances, the cost of searching the entire neighborhood every iteration becomes prohibitive so that SMA-S and CLS solve no instances. It is immediately apparent

TABLE I
NUMBER OF RUNS (OUT OF 250) IN WHICH A SOLUTION WAS IDENTIFIED FOR DIFFERENT-LENGTH MAX-3SAT PROBLEMS

Algorithm	Length 50	Length 100
GA	125	21
SMA-S	154	0
SMA-G	153	25
CLS	141	0
CLS-I	141	3
CLG	135	21
CLG-I	136	25
CLA	144	8
CLA-I	142	8
CTS	112	10
CTS-I	131	12
CTG	117	19
CTG-I	118	20
CTA	119	14
CTA-I	126	19

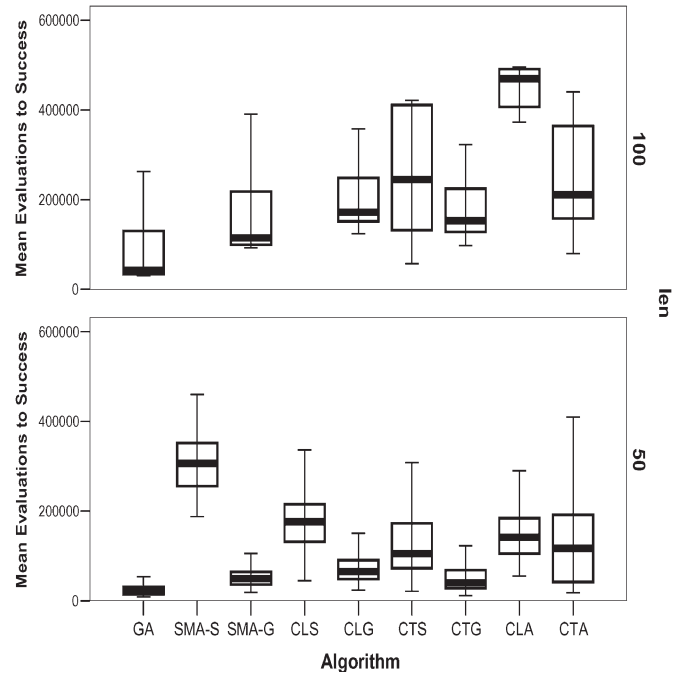


Fig. 2. Box plots of average evaluations to success for ten runs on each of 25 MAX-3SAT instances with (top) 100 and (bottom) 50 variables.

that the adaptive variants CLA and CTA appear to perform on a par with whichever of the S or G variants is better for each length, suggesting successful adaptation.

The effect of the ability to use inversion is less clear cut, but it yields a slight improvement on the longer instances.

C. Efficiency

Fig. 2 shows the mean time to solution analyzed by algorithm and problem size. ANOVA on just the COMA results showed that the ability to use inversion was not a significant factor on the AES, so these results are omitted from the figure for clarity. As can be seen, the GA is the fastest algorithm followed by a close grouping of SMA-G, CLG, and CTG, with the CLS and CTS variants taking more time and having a higher variance. The adaptive pivot variants both fall between their




Fig. 3. Box plots of best values found for ten runs on each of 25 MAX-3SAT instances with 250 variables.

respective greedy and steepest counterparts both in terms of mean and variance. The fitness-based pairing is faster than the linked counterpart in every case. A two-way ANOVA with problem size and algorithm as fixed factors showed that the algorithm was a determining factor in AES. Post hoc testing using Tamhane’s test at the 95% confidence level showed that the GA was significantly faster and the SMA-S significantly slower than the other algorithms, but otherwise, the grouping was not well defined, with most algorithms being not significantly different to three or more others.

D. MBF

Fig. 3 shows the MBF analyzed by algorithm for the longer instances. As noted above, the performance of the steepest ascent versions of any of the algorithms was worse than the greedy or adaptive versions; therefore, these are omitted for clarity. Analysis shows no significant difference between the means for linked and fitness-based pairings (CTG–CLG, CTA–CLA, etc.) for this performance indicator, although the fitness-based pairing algorithms tend to have higher variance than their linked pairing counterparts. Again, the use of inversion was not a significant factor for these results.

Overall, adding the possibility of an “invert” symbol (#) to the action has clearly improved the SR and has not been detrimental to the other measures. The variants with adaptive pivot rules are outperformed by either the steepest or the greedy ascent, whichever is better, but this is often marginal. In particular, if the most important criterion to the algorithm designer is reliability (high SR), then letting the pivot rules evolve is clearly preferable to making a fixed choice. On some problems,

steepest was better, on others, greedy, and on the MAX-3SAT problems, the choice is size dependent. In contrast to this, if the desired properties are speed (low AES) and consistently good results (high MBF), then the fixed greedy strategy is preferable on the MAX-3SAT, although, again, this does not always hold for other problems.

VII. CONCLUSION

The premise that evolutionary optimization algorithms can be improved by incorporating an appropriate local search mechanism is now widely accepted, but it is increasingly recognized that the choice of move operator and, hence, neighborhood function used in the LS is crucial to delivering success.

This paper reviewed a framework that offers the promise for creating robust scalable optimization techniques based on the concept of coevolving memes encoding definitions of LS operators. Results showing the development of algorithms fitting into this framework with progressively more coevolutionary natures and more degrees of freedom available to the evolution of memes were reviewed. These results are highly competitive on a variety of different classes of problem. These experiments have used a number of different performance metrics to compare between algorithms, as different intended uses of EAs have different goals. In some repeated tasks, reliably fast and reasonably accurate behavior (i.e., low AES and high MBF) may be more important than finding the true global optimum at whatever cost (e.g., high SR and AES), but the reverse may be true in, e.g., design situations.

The results obtained illustrate that performance improvement can arise from different mechanisms. If the representation of the rules is able to capture regular repeated features within the problem space, then highly scalable behavior is exhibited—for example, the linear speedup on the 4-Trap function. This arises from the rapid evolution of the system to a rule set that captures and represents knowledge about how to solve the problem. It was noted that in order for this to occur, it is necessary to maintain sufficient accurate selection pressure within the population of local searchers.

In contrast to this, when there is no sufficient selective pressure for evolution, a “fallback” position is observed. This might occur, for example, when a “good” pattern only applies to one position in the solution, or when the rule representation cannot possibly capture the regularities present in the space. In these cases, improved reliability is observed, but at the expense of speed of solution. Essentially, what happens is that even if the solution population converges to a local optimum for its crossover and mutation operators, the continued evolution of memes means that, eventually, an LS landscape is discovered in which the solutions are not locally optimal and improvements can occur. This is akin to variable neighborhood search or other adaptive MAs, but with the advantage that it is not necessary to specify, or be bound by, a fixed set of neighborhood functions.

These results from a series of benchmark tests highlight this problem of choosing the appropriate LS operator that provided the original rationale for the development of COMA. For example, although the MA with a simple bit-flipping hill climber had the highest SRs and MBF on the MAX-3SAT problems,

its performance on the other problems was derisory and frequently worse than the simple GA.

The COMA variants using “linked” pairing, which effectively self-adapt the memes, exhibited better performance than the GA or SMA over a wide range of problems according to different metrics. The results on MAX-3SAT were comparable with SMA, but this problem is perhaps unusual in having no structure to be exploited. Other authors have suggested methods for reordering MAX-SAT representations so as to maximize local gene interactions, which would clearly aid the rule-based COMA algorithms, as probably would by allowing more evaluations to try different neighborhoods.

When the behavior of the truly coevolutionary models with fitness-based selection of memes (CTS and CTG) is considered, the picture is less clear. On the H-IFF and multiple 4-Trap problems, the method works well, particularly when used with the steepest ascent pivot rule. However, on the MAX-3SAT problem, the situation is reversed, and the greedy ascent rule appears to work better. As suggested, this may be because the neighborhoods are so large or because the fitness measure used for memes was too simple. The use of adaptive pivot rules goes a long way toward resolving the issue of choosing a pivot rule, but there is clearly scope for future work here, and it is worth considering a change from what Ong *et al.* would term local to global fitness measures such as running averages, etc.

Clearly, there remains much work to be done in analyzing the possibilities of this framework. It would be fatuous to claim that COMA represents some fabulous all-purpose problem solver, however promising these results. Immediate priorities are the investigation of alternative methods for credit assignment within the meme population and the extension to more generic representations of conditions and actions. However, the two modes of operation noted above, coupled with the ability of the algorithm to explicitly represent the information that it has learned and is using to solve the problem at hand, would seem to offer much potential.

ACKNOWLEDGMENT

The author would like to thank N. Krasnogor for many fruitful discussions during the initial stages of this work.

REFERENCES

- [1] N. Krasnogor and J. Smith, “A tutorial for competent memetic algorithms: Model, taxonomy and design issues,” *IEEE Trans. Evol. Comput.*, vol. 9, no. 5, pp. 474–488, Oct. 2005.
- [2] W. Hart, N. Krasnogor, J. Smith, Eds. *Recent Advances in Memetic Algorithms*. Berlin, Germany: Springer-Verlag, 2004.
- [3] P. Moscato. (2005, Jul.). *Memetic Algorithms*. [Online]. Available: http://www.densis.fee.unicamp.br/~moscato/memetic_home.html
- [4] J. Smith, “Co-evolution of memetic algorithms: Initial investigations,” in *Proc. 7th Conf. Parallel Problem Solving From Nature*, J. M. Guervos, P. Adamidis, H.-G. Beyer, J.-L. Fernandez-Villacanas, and H.-P. Schwefel, Eds. Berlin, Germany: Springer-Verlag, 2002, vol. 2439, pp. 537–548.
- [5] —, “Protein structure prediction with co-evolving memetic algorithms,” in *Proc. CEC*, Piscataway, NJ, 2003, pp. 2346–2353.
- [6] —, “The co-evolution of memetic algorithms for protein structure prediction,” in *Recent Advances in Memetic Algorithms*, W. Hart, N. Krasnogor, and J. Smith, Eds. Berlin, Germany: Springer-Verlag, 2004, pp. 105–128.
- [7] —, “Co-evolving memetic algorithms: A learning approach to robust scalable optimisation,” in *Proc. CEC*, Piscataway, NJ, 2003, pp. 498–505.
- [8] Y. Ong, M. Lim, N. Zhu, and K. Wong, “Classification of adaptive memetic algorithms: A comparative study,” *IEEE Trans. Syst., Man, Cybern. B, Cybern.*, vol. 36, no. 1, pp. 141–152, Feb. 2006.
- [9] N. Krasnogor, “Coevolution of genes and memes in memetic algorithms,” in *Proc. Genetic and Evol. Comput. Conf. Workshop Program*, A. Wu, Ed., 1999.
- [10] N. Krasnogor and J. Smith, “A memetic algorithm with self-adaptive local search: TSP as a case study,” in *Proc. GECCO*, D. Whitley, D. Goldberg, E. Cantu-Paz, L. Spector, I. Parmee, and H.-G. Beyer, Eds. San Francisco, CA, 2000, pp. 987–994.
- [11] —, “Emergence of profitable search strategies based on a simple inheritance mechanism,” in *Proc. GECCO*, L. Spector, E. Goodman, A. Wu, W. Langdon, H.-M. Voigt, M. Gen, S. Sen, M. Dorigo, S. Pezeshek, M. Garzon, and E. Burke, Eds. San Francisco, CA, 2001, pp. 432–439.
- [12] N. Krasnogor, “Studies in the theory and design space of memetic algorithms,” Ph.D. dissertation, Univ. West England, Bristol, U.K., 2002.
- [13] Y. Ong and A. Keane, “Meta-Lamarckian learning in memetic algorithms,” *IEEE Trans. Evol. Comput.*, vol. 8, no. 2, pp. 99–110, Apr. 2004.
- [14] P. Cowling, G. Kendall, and E. Soubeiga, “A hyperheuristic approach to scheduling a sales summit,” in *Lecture Notes in Computer Science*, vol. 2079. Berlin, Germany: Springer-Verlag, 2001, pp. 176–195.
- [15] E. Burke and A. Smith, “Hybrid evolutionary techniques for the maintenance scheduling problem,” *IEEE Trans. Power Syst.*, vol. 15, no. 1, pp. 122–128, Feb. 2000.
- [16] G. Kendall, P. Cowling, and E. Soubeiga, “Choice function and random hyperheuristics,” in *Proc. 4th Asia-Pac. Conf. SEAL*, 2002, pp. 667–671.
- [17] E. Burke, G. Kendall, and E. Soubeiga, “A tabu search hyperheuristic for timetabling and rostering,” *J. Heuristics*, vol. 9, no. 6, pp. 451–470, Dec. 2003.
- [18] N. Krasnogor, “Self-generating metaheuristics in bioinformatics: The protein structure comparison case,” in *Genetic Program. Evolvable Mach.*, vol. 5, Jun. 2004, pp. 181–201.
- [19] N. Krasnogor and S. Gustafson, “A study on the use of ‘self-generation’ in memetic algorithms,” *Nat. Comput.*, vol. 3, no. 1, pp. 53–76, 2004.
- [20] P. Hansen and N. Mladenović, “An introduction to variable neighborhood search,” in *Proc. MIC Conf. Meta-Heuristics: Adv. and Trends Local Search Paradig. Optim.*, S. Voß, S. Martello, I. Osman, and C. Roucairol, Eds. Dordrecht, The Netherlands, 1998, pp. 433–458.
- [21] J. Smith and T. Fogarty, “Operator and parameter adaptation in genetic algorithms,” *Soft Comput.*, vol. 1, no. 2, pp. 81–87, 1997.
- [22] R. Hinterding, Z. Michalewicz, and A. Eiben, “Adaptation in evolutionary computation: A survey,” in *Proc. IEEE Conf. Evol. Comput.*, Piscataway, NJ, 1997, pp. 65–69.
- [23] A. Eiben, R. Hinterding, and Z. Michalewicz, “Parameter control in evolutionary algorithms,” *IEEE Trans. Evol. Comput.*, vol. 3, no. 2, pp. 124–141, Jul. 1999.
- [24] H.-P. Schwefel, *Numerical Optimisation of Computer Models*. New York: Wiley, 1981.
- [25] D. Fogel, “Evolving artificial intelligence,” Ph.D. dissertation, Univ. California, San Diego, 1992.
- [26] T. Bäck, “Self adaptation in genetic algorithms,” in *Proc. 1st Eur. Conf. Artif. Life—Toward a Practice of Autonomous Systems*, F. Varela and P. Bourgine, Eds. Cambridge, MA, 1992, pp. 263–271.
- [27] J. Smith and T. Fogarty, “Self adaptation of mutation rates in a steady state genetic algorithm,” in *Proc. IEEE Conf. Evol. Comput.*, Piscataway, NJ, 1996, pp. 318–323.
- [28] J. Schaffer and A. Morishima, “An adaptive crossover distribution mechanism for genetic algorithms,” in *Proc. 2nd Int. Conf. Genetic Algorithms and Their Appl.*, J. Grefenstette, Ed. Hillsdale, NJ, 1987, pp. 36–40.
- [29] J. Smith and T. Fogarty, “Adaptively parameterised evolutionary systems: Self adaptive recombination and mutation in a genetic algorithm,” in *Proc. 4th Conf. Parallel Problem Solving From Nature*, H.-M. Voigt, W. Ebeling, I. Rechenberg, and H.-P. Schwefel, Eds. Berlin, Germany: Springer-Verlag, 1996, vol. 1141, pp. 441–450.
- [30] J. Paredis, “The symbiotic evolution of solutions and their representations,” in *Proc. 6th Int. Conf. Genetic Algorithms*, L. Eshelman, Ed. San Francisco, CA, 1995, pp. 359–365.
- [31] M. Potter and K. De Jong, “A cooperative coevolutionary approach to function optimisation,” in *Proc. 3rd Conf. Parallel Problem Solving From Nature*, Y. Davidor, H.-P. Schwefel, and R. Männer, Eds. Berlin, Germany: Springer-Verlag, 1994, vol. 866, pp. 248–257.
- [32] L. Bull, “Artificial symbiology,” Ph.D. dissertation, Univ. West England, Bristol, U.K., 1995.
- [33] S. Kauffman, *Origins of Order: Self-Organization and Selection in Evolution*. New York: Oxford Univ. Press, 1993.

- [34] L. Bull and T. Fogarty, "Horizontal gene transfer in endosymbiosis," in *Proc. 5th Int. Workshop Artif. Life: Synthesis and Simul. Living Syst. (ALIFE)*, C. Langton and K. Shimohara, Eds. Cambridge, MA, 1997, pp. 77–84.
- [35] L. Bull, "Evolutionary computing in multi agent environments: Partners," in *Proc. 7th Int. Conf. Genetic Algorithms*, T. Bäck, Ed. San Francisco, CA, 1997, pp. 370–377.
- [36] G. Parker and H. Blumenthal, "Varying sample sizes for the co-evolution of heterogeneous agents," in *Proc. CEC*, 2004, pp. 766–771.
- [37] J. Paredis, "Coevolutionary algorithms," in *Handbook of Evolutionary Computation*, T. Bäck, D. Fogel, and Z. Michalewicz, Eds. New York: Oxford Univ. Press, 1998.
- [38] S. Luke and L. Spector, "Evolving teamwork and coordination with genetic programming," in *Proc. 1st Annu. Conf. Genetic Program.*, J. Koza, D. Goldberg, D. Fogel, and R. Riolo, Eds. Cambridge, MA, 1996, pp. 141–149.
- [39] H. Kargupta and S. Ghosh, "Towards machine learning through genetic code-like transformations," *Comput. Sci. Electr. Eng. Dept.*, Univ. Maryland Baltimore County, Baltimore, Tech. Rep. TR-CS-01-10, 2001.
- [40] R. Keller and W. Banzhaf, "Genetic programming using genotype-phenotype mapping from linear genomes into linear phenotypes," in *Proc. 1st Annu. Conf. Genetic Program.*, J. Koza, D. Goldberg, D. Fogel, and R. Riolo, Eds. Cambridge, MA, 1996, pp. 116–122.
- [41] —, "The evolution of genetic code in genetic programming," in *Proc. GECCO*, W. Banzhaf, J. Daida, A. Eiben, M. Garzon, V. Honavar, M. Jakiela, and R. Smith, Eds. San Francisco, CA, 1999, pp. 1077–1082.
- [42] W. Hart, "Adaptive global optimization with local search," Ph.D. dissertation, Univ. California, San Diego, 1994.
- [43] T. Jones, "Evolutionary algorithms, fitness landscapes and search," Ph.D. dissertation, Univ. New Mexico, Albuquerque, 1995.
- [44] P. Merz and B. Freisleben, "Fitness landscapes and memetic algorithm design," in *New Ideas in Optimization*, D. Corne, M. Dorigo, and F. Glover, Eds. London, U.K.: McGraw-Hill, 1999, pp. 245–260.
- [45] T. Bäck, D. Fogel, and Z. Michalewicz, Eds., *Handbook of Evolutionary Computation*. New York: Oxford Univ. Press, 1997.
- [46] K. Dill, "Theory for the folding and stability of globular proteins," *Biochemistry*, vol. 24, no. 6, pp. 1501–1509, Mar. 1985.
- [47] B. Berger and T. Leight, "Protein folding in the hydrophobic-hydrophilic (hp) model is NP-complete," in *Proc. 2nd Annu. Int. Conf. Comput. Molecular Biol. RECOMB*, 1998, pp. 30–39.
- [48] R. Unger and J. Moult, "Genetic algorithms for protein folding simulations," *J. Theor. Biol.*, vol. 231, no. 1, pp. 75–81, 1993.
- [49] N. Krasnogor, W. Hart, J. Smith, and D. Pelta, "Protein structure prediction with evolutionary algorithms," in *Proc. GECCO*, W. Banzhaf, J. Daida, A. Eiben, M. Garzon, V. Honavar, M. Jakiela, and R. Smith, Eds. San Francisco, CA, 1999, pp. 1596–1601.
- [50] Satlib—The satisfiability library. [Online]. Available: <http://www.satlib.org>



Jim E. Smith took the Electrical Sciences Tripos at Cambridge University, Cambridge, U.K., and then worked in industry for several years before returning to study. He received both the M.Sc. degree (with distinction) in communicating computer systems and the Ph.D. degree from the University of the West of England, Bristol, U.K., in 1993 and 1998, respectively.

He has been with the University of the West of England since 1993, first as a Research Fellow in adaptive systems, and currently as a Senior Lecturer

in the School of Computer Science. He has been researching in the field of heuristic optimization, machine learning, and data mining since 1994 and has been awarded a number of industrially, U.K. governmental agency-, and European Union-funded projects. He has authored a book on evolutionary computing and has recently edited a book and a special issue of *Evolutionary Computation on Memetic Algorithms* (a class of hybrid search heuristics).

Dr. Smith is a Cofounder and the Chair of the International Workshops on Memetic Algorithms and was a Co-Chair of the 8th International Conference on Parallel Problem Solving from Nature 2004.